

# Can You Polish Your Dutch?

Jacek Pardyak

November 8th, 2017

# Outline

- ▶ Business understanding
- ▶ Data understanding
- ▶ Data preparation
- ▶ Model construction
- ▶ Model evaluation
- ▶ Insights from the data

# Business understanding

To communicate people use words. Words are composed of letters over an alphabet.


Letters common for the two languages: a, b, c, d, e, f, g, h, i, j, k, l, m, n, o, p, r, s, t, u, w, y, z.

Letters exclusively used in Polish: ą, ć, ę, ł, ń, ó, ś, ź, ż

Letters exclusively used in Dutch: q, v, x

## Problem understanding

Build a model which distinguishes Polish words from Dutch.



Eventually find (dis)similarities of the two languages.

# Data understanding

To train and test models we use **Aspell** dictionaries.

Sample of the Polish data:

---

nieprzegadywanie/UV  
snowboard/NQsT  
najkompletniej  
Komiaczka/MmN  
synonimowy/bXxY  
ziębał

---

Polish is inflected language. Symbols after / are used to mark affixes.

## Data understanding cont.

Sample of the Dutch data:

---

aanbrei  
lichtwaterreactor  
ouderhuis  
inhoudstafels  
selectiekamer  
linkerhelft

---

We use 341461 Dutch and 289840 Polish words. After applying 7k inflection rules number of words in Polish grows to 3.8M.

# Data preparation

1. Remove everything after “/” sign

```
words$Word <- gsub("\\\\/.*", "", words$Word)
```

2. Split words into list of letters

```
words$Word_split <- lapply(words$Word, function(x) {  
  paste(unlist(strsplit(x, "")), collapse = " ")})
```

3. Coerce list of words into “document-term-matrix”

```
dtm <- DocumentTermMatrix(Corpus(VectorSource(  
  words$Word_split)), control = list(  
    tokenize = UnicodeTokenizer, wordLengths = c(1,2)))
```

## Data preparation cont.

For example words:

```
## [1] "Adrianna" "Rea"
```

are represented as:

```
## <<DocumentTermMatrix (documents: 2, terms: 40)>>
## Non-/sparse entries: 8/72
## Sparsity           : 90%
## Maximal term length: 2
## Weighting          : term frequency (tf)
## Sample            :
##           Terms
## Docs      a b c d e i n p r s
##   311      3 0 0 1 0 1 2 0 1 0
##  31197  1 0 0 0 1 0 0 0 1 0
```



## Data preparation cont.

4. Create output variable of two classes (1 for Dutch word, 0 otherwise)
5. Split the data into training and test data sets (70/30)

Finally:

- ▶ `X_train` - matrix of 441910 rows and 40 columns stores input variables of training data set
- ▶ `Y_train` - vector of 441910 elements stores output variable of training data set
- ▶ `X_test` - matrix of 189391 rows and 40 columns stores input variables of test data set
- ▶ `Y_test` - vector of 189391 elements stores output variable of test data set

# Model construction

To construct our first Deep Neural Network model we need to perform following steps:

- ▶ initialize the model,
- ▶ add layers to the model,
- ▶ compile and fit our model.

```
# Load 'keras' - API to 'TensorFlow' engine  
require(keras)  
# Apply one-hot-bit encoding  
Y_train <- to_categorical(Y_train)  
# Construct an empty sequential model  
# composed of a linear stack of layers  
model <- keras_model_sequential()
```

## Model construction cont.

```
model %>%  
  # add a dense layer  
  layer_dense(units = 500, input_shape = 40,  
              kernel_initializer="glorot_uniform",  
              activation="sigmoid") %>%  
  # add dropout to prevent overfitting  
  layer_dropout(rate = 0.5) %>%  
  layer_dense(units = 300,  
              kernel_initializer="glorot_uniform",  
              activation="sigmoid") %>%  
  layer_dropout(rate = 0.5) %>%  
  layer_dense(units = 100,  
              kernel_initializer="glorot_uniform",  
              activation="sigmoid") %>%  
  layer_dropout(rate = 0.5) %>%  
  layer_dense(units = 2,  
              kernel_initializer="glorot_uniform",  
              activation="softmax")
```

## Model construction cont.

```
# Compile the model
```

```
model %>%
```

```
  compile(loss = 'categorical_crossentropy',  
          optimizer = optimizer_adam(lr=0.001,  
                                     beta_1=0.9,  
                                     beta_2=0.999,  
                                     epsilon=1e-08,  
                                     decay=0.0),  
          metrics = 'accuracy')
```

```
# Fit the model from the training data
```

```
training_history <- model %>%
```


```
# batch_size - number of samples per gradient update
```

```
# epochs - number of times to iterate on a dataset
```

```
fit(X_train, Y_train, batch_size = 64,  
    epochs = 12, verbose = 1,  
    validation_split = 0.1)
```

# Model construction cont.

Training history:



## Model evaluation

```
# Make predictions on the test dataset
```

```
Y_test_hat <- model %>%  
  predict_classes(X_test)
```

```
Y_test_hat <- as.integer(Y_test_hat)
```

```
##           Y_test_hat  
## Y_test      0      1  
##      0 78309  8592  
##      1  8164 94326
```

```
## accuracy:  91.15269%  
## precision:  91.65161%  
## recall:     92.03434%  
## f-measure:  91.84258%
```

# Results


We applied Neural Networks to identify language of a word just using letter frequencies.

Performance of the model is very good - accuracy, precision and recall above 90%.

Further we try to understand this behavior.

# Insights from data

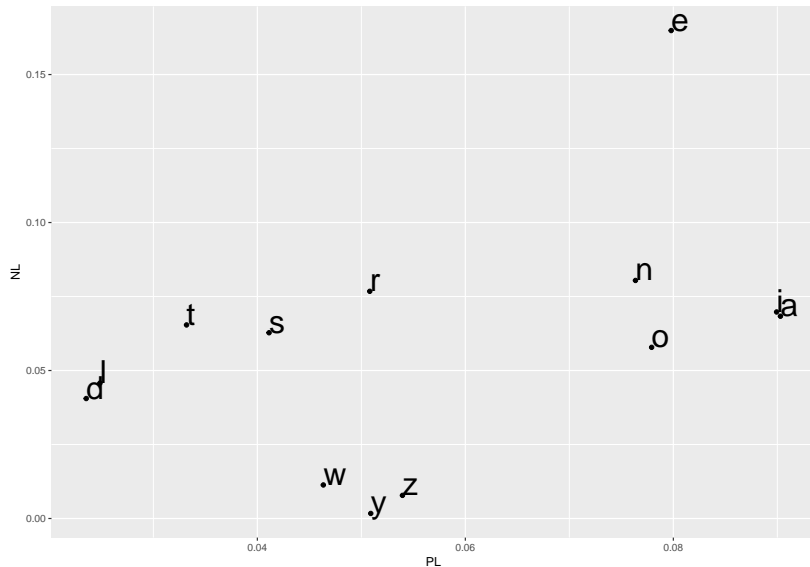
Word length comparison:





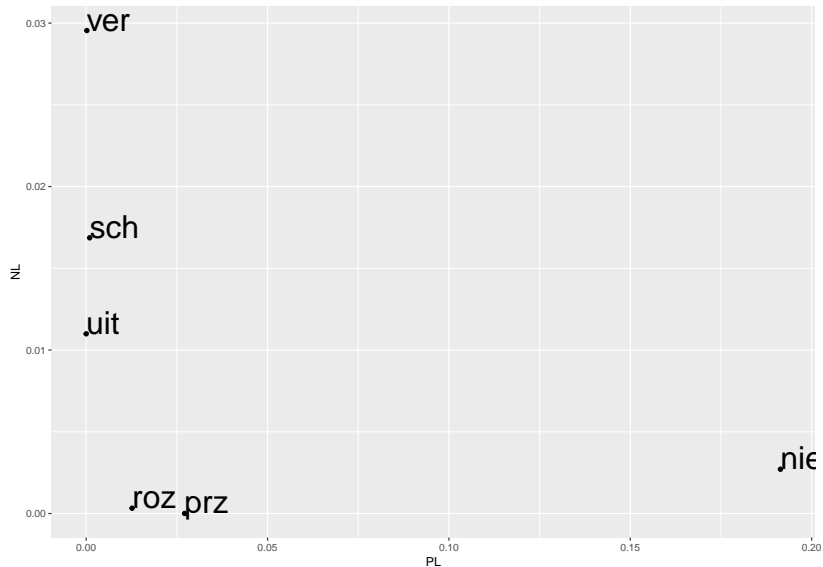
# Insights from data cont.

Relative letter frequency:



## Insights from data cont.

Relative frequency of initial trigrams:



## Insights from data cont.

Polish words perfectly and approximately matching Dutch words:


- ▶ ananas, balkon, chaos, duet, echo, filet, gratis, handel, impotent, jacht, kapsel, legenda, wiek and 3.3k more
- ▶ abiturient ~ abituriënt, banan ~ banaan, bestseler ~ bestseller, dermatolog ~ dermatoloog, fortepian ~ fortepiano, wachta ~ wacht and 2.6k more

Our DNN model was trained on words assigned to two different classes.

Watch 'false friends' - words spelled the same but meaning something different.

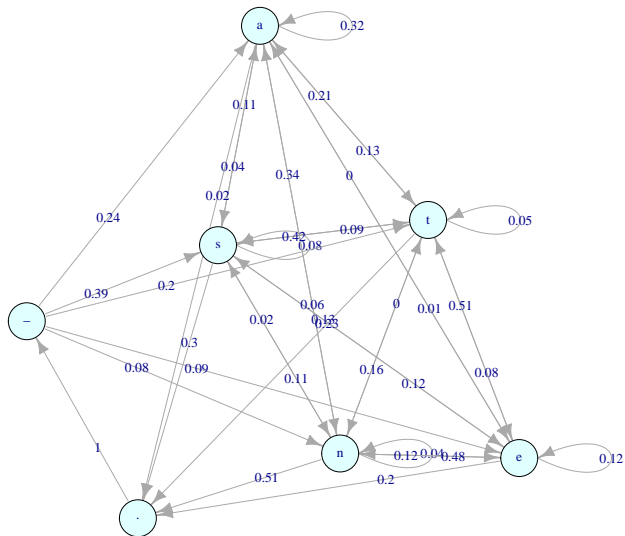
## Insights from data cont.

We can use Markov chains to build probabilistic model of a language. Excerpt of the Polish model:



## Insights from data cont.

Excerpt of the Dutch language probabilistic model:



## Insights from data cont.

Animated most likely “random” walk through the Dutch graph:

## Insights from data cont.

Probabilistic language models can be used to generate 'synthetic' words:

- ▶ wypcy, ośm, donie, bonijny, tać, nionwry, szero, zberemy
- ▶ vevon, orin, veden, gaaauk, ilin, ommouin, pamoe, parle

Our model accurately recognized language of these synthetic words.

# Summary

- ▶ Deep Learning is a very powerful technique
- ▶ Use of bi- and trigrams will lead to even better performance
- ▶ Dutch and Polish are dissimilar languages
- ▶ About 3% of words is commonly used in Polish and Dutch